

Name: Emil Ljung
Institution: Computer science
Educational program: Game programming
Course: DV1463 Performance optimization
Date: 2017-05-07

Project 2 - Multithreaded Programming using Pthreads

Note: All the files which were used for the assignments are included in the zip-file and are named so that it's easy to identify which file was used for which task. For instance, matmul_8threads.c was used for getting the execution time for blocked matrix multiplication with 8 threads.
All files are compiled with: `gcc -std=c99 {filename} -lm -lpthread -o {output name}`
Run the code with: `/usr/bin/time ./{output name}`

Fractal

Sequential: 3.21 seconds
1 thread: 3.12 seconds
8 threads: 1.32 seconds

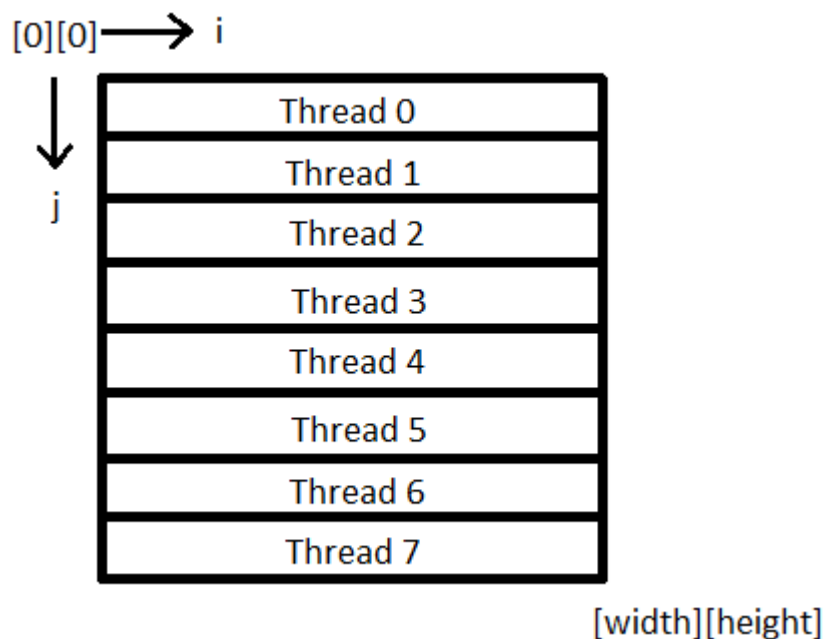
Speedup: 243.2%

The logic behind splitting up the work between the threads in the files `fractal_1thread.c` & `fractal_8threads.c` are the same. The only exception is that the `#define` variable `NR_OF_THREADS` equals 8 in `fractal_8threads.c` and 1 in `fractal_1thread.c`. All the threads which are going to be created exists in an array called `threads`. Each thread have a struct each which belongs to them (all the structs can be found in the array `t_data`). The struct is simply called `thread_data` and consists of 1 int pointer for the pixel map and 2 int values which tells from and to which `j` the threads shall work on. The 2 `j` values are based on `NR_OF_THREADS` and the `#define` variable `width`, which equals 1024. This way we can get a parallelization. In the figure below you can see how I imagined it was done. However, it is NOT what is actually happening in the code.

All threads are created in the function `init_threads_mandelbrot`. This is also where the structs get their values which will be used in the function `mandelbrot`, which is the function each thread will run. There you'll see the 2 for loops which are:

```
for (i = 0; i < width; i++) {  
    for(j = data->start_j; j < data->end_j; j++) {  
        ...  
    }  
}
```

This code snippet is what made me treat the parallelization as the figure below shows since it made me think about a 2D array and thus treat it as a such.



How I imagined the partitioning between each thread.

Blocked Matrix-Matrix multiplication

Row wise 8 threads: 11.69 seconds


Blocked 8 threads: 13.57 seconds

The way I solved this is similar to how the fractal was threaded. Basically, the difference is that in this solution the `thread_data` struct also have `int i_start & i_end`, which are used like the `j` counterparts found in the struct in the fractal assignment.

Since there was no requirement to implement a 1 thread version of this I only focused on the solution for 8 threads. There were also nothing that stated exactly how the partitioning was supposed to be between the threads, so I split up the work the way the figure below shows.

The code for splitting up the work looks like this:

```
for(i = 0; i < NR_OF_THREADS*0.5; i++) {  
    int k = i*2;  
    t_data[k].start_i = 0;  
    t_data[k].end_i = SIZE*0.5;  
    t_data[k].start_j = SIZE/(NR_OF_THREADS*0.5) * i;  
    t_data[k].end_j = SIZE/(NR_OF_THREADS*0.5) * (i+1);  
    int t0 = pthread_create(&threads[k], NULL, row_calculation_blocked, &t_data[k]);  
  
    t_data[k+1].start_i = SIZE*0.5;  
    t_data[k+1].end_i = SIZE;  
    t_data[k+1].start_j = t_data[k].start_j;  
    t_data[k+1].end_j = t_data[k].end_j;  
    int t1 = pthread_create(&threads[k+1], NULL, row_calculation_blocked, &t_data[k+1]);  
}
```



Quick Sort

Sequential: 30.61 seconds
1 thread: 31.53 seconds
8 threads: 11.63 seconds

Speedup: 263.2%

For this one I specifically went for the 8 threads solution first. What I did here was that I first created the root thread which pretty much just runs `thread_quick_sort`. This function is in several ways like the function `quick_sort` but when `quick_sort` isn't called a new thread which runs `thread_quick_sort` is called. After the parent thread has created its 2 child threads (1 for the values to the left of the pivot and 1 for the ones to the right in the array) it'll call `pthread_join` for both child threads, thus making the parent waiting for both of them to complete their tasks. This means that if I set `NR_OF_THREADS` to 8, which at first seems quite reasonable, then, if we look at the figure below, we'll see that the tree will have 4 leaf threads (0-7 are created). Since the parent threads are just waiting for the children we only have 4 ACTIVE threads. In order to create 8 active leaf threads `NR_OF_THREADS` must be set to 15!

Since I, in `thread_quick_sort`, have the if-statement `if(childID > NR_OF_THREADS-1)` which starts sorting if it's true I need to make sure that I create threads with unique IDs. Simply so I don't try to create a thread which already have been created. The `thread_data` consists of `thread_number`, `low` & `high` (all of which are int). The variable `thread_number` is used to set the child IDs, which are set to $\text{thread_number} * 2 + 1$ for the child thread created to the left of pivot and $\text{thread_number} * 2 + 2$ for the other child thread. By creating the threads and giving them $\text{thread_number} * 2 + 1$ or $+ 2$ I always end up with the same tree shown in the figure below and with the children at the same place!

This way of creating each child thread only works for binary trees and creates a special case if only 1 thread is to be created (unless some logic is to be changed).

If you would use the code which was used for 8 threads and then just set `NR_OF_THREADS` to 1 then 2 threads would be created in total anyway. In order to use 1 thread in a simple way I instead set `NR_OF_THREADS = 0` and both the struct & thread array have the size 1 (instead of using the `NR_OF_THREADS` to set their size). This way only the root thread is created and in `thread_quick_sort` it goes directly into the `quick_sort` functions and starts sorting.

